PRAC4 : Matrix Multiplication

DIETZ T., JABER A., DONNENFELD T.

November 17, 2025

1 Step 1: Theoretical Analysis

1.1 Iterative Version: Work Partitioning

For the classical iterative matrix multiplication, we parallelize the computation by distributing the **rows of the output matrix** across the available threads. Matrix multiplication has no data dependencies between rows of C because:
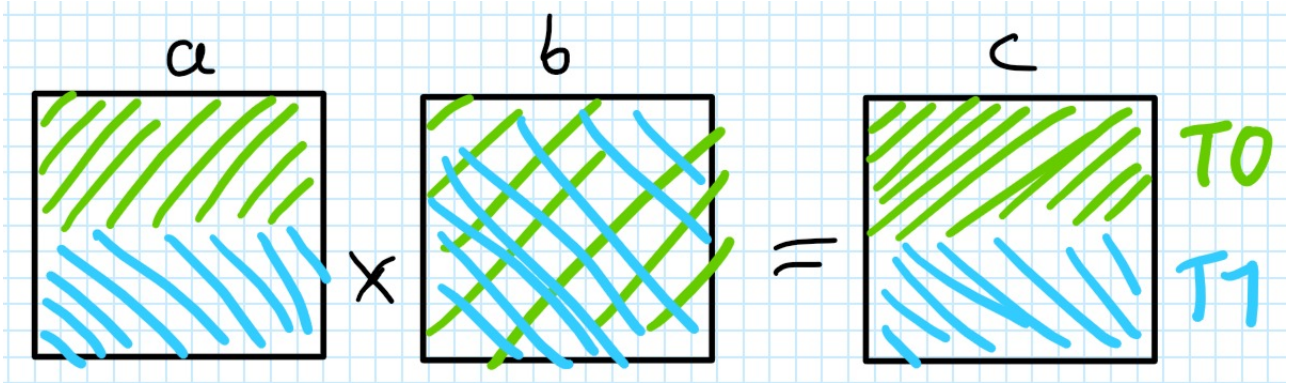
$$C[i,j] = \sum_{k=0}^{N-1} A[i,k]\,B[k,j],$$

With k threads, we partition rows of C into k blocks

Each thread:

- reads $\frac{N}{k}$ rows of A,
- reads all of B,
- writes to disjoint rows of C.

This gives perfect load balance if matrix has data evenly distributed across its rows and columns and no race conditions.

In particular for two threads:

1.2 Divide and Conquer Version

The recursive version repeatedly divides the problem into 8 independent subproblems of size $\frac{SZ}{2} \times \frac{SZ}{2}$ until the base-case size $DQSZ$ is reached.

By analogy with the same 1 dimension problem : we find the number of required splits steps $n = \log_2\left(\frac{N}{DQSZ}\right)$

At recursion level i, the number of tasks is 8^i

The total number of tasks is the sum of the terms of the previous geometric sequence, n being the max number of splits seen above. Therefore : $T = \sum_{i=1}^{n} 8^i = \frac{8^{n+1}-8}{7}$.

The maximum concurrency is the number of "leaf" tasks when all tasks have been created: 8^s

Case for $N = 1024$, $DQSZ = 512$

Here $s = \log_2(1024/512) = 1$, so only one level of recursion.

$$s = 1, \qquad T = 8, \qquad Max\,concurrency = 8.$$

Memory Access per Task

We have to be careful with the tasks : sums and products of A and B submatrices can run in parallel, but these results have to be added two by two to form submatrices of C : this causes values in C to be accessed by two tasks possibly at the same time.

2 Step 2: Practical Parallel Implementation

2.1 Iterative Version

We change loop order to maximize cache re-use and we unroll the loops to take advantage of contiguous memory operations and potential SIMD vectorization.

```
#pragma omp parallel for
for (int i=0; i<N; i+=2)
  for (int k=0; k<N; k+=2)
    for (int j=0; j<N; j++)
    {
      type B1 = b[k*N+j];
      type B2 = b[(k+1)*N+j];

      c[i*N+j]     += a[i*N+k]     *B1 + a[i*N+k+1]     *B2;
      c[(i+1)*N+j] += a[(i+1)*N+k] *B1 + a[(i+1)*N+k+1] *B2;
    }
```

This reduces memory accesses.

2.2 Divide and Conquer Version

We parallelize the recursion using OpenMP tasks:

```
#pragma omp parallel
{
  #pragma omp single
  {
    #pragma omp task { ... }
    #pragma omp task { ... }
    #pragma omp task { ... }
    #pragma omp task { ... }
  }
}
```

3 Step 3: Performance Analysis

We benchmarked $N = 2048$ for a range of thread counts and $DQSZ$ values. Counters collected using `perf stat`:

- execution time,

- instructions, cycles, IPC,

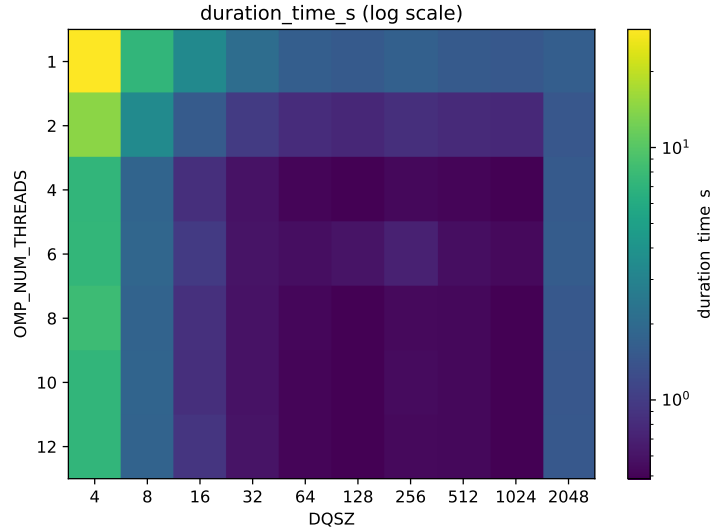- cache misses,

- speedup and efficiency.

We established the following single-threaded baseline execution time :

dq	duration (s)
2	195.03
4	29.36
8	7.23
16	3.42
32	2.10
64	1.65
128	1.56
256	1.72
512	1.55
1024	1.51
2048	1.66

Table 1: Single-thread execution time as a function of dq

We then performed analysis using all optimizations provided.
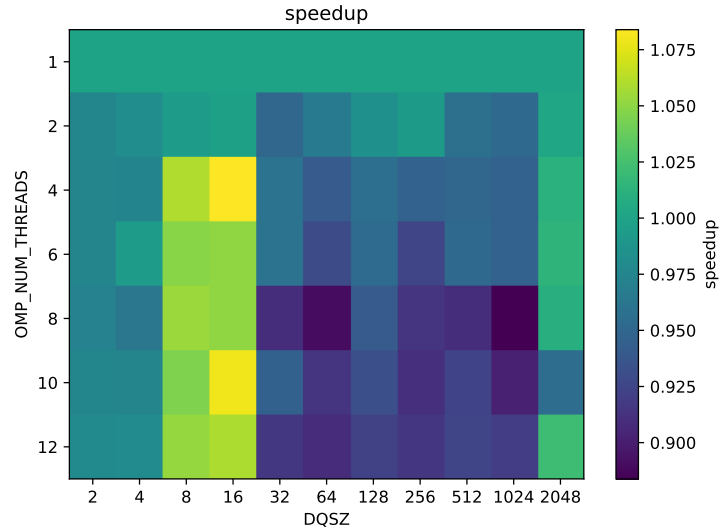
3.1 Heatmap 1: Execution Time

Observations:

- Large and very small $DQSZ$ (4, 8, 2048) yields very poor performance

- Single threaded runs yield very poor performance

- $DQSZ = 64$ or $DQSZ = 128$ produce the best runtimes.

This graph shows the overall best configurations for this problem on the machine.

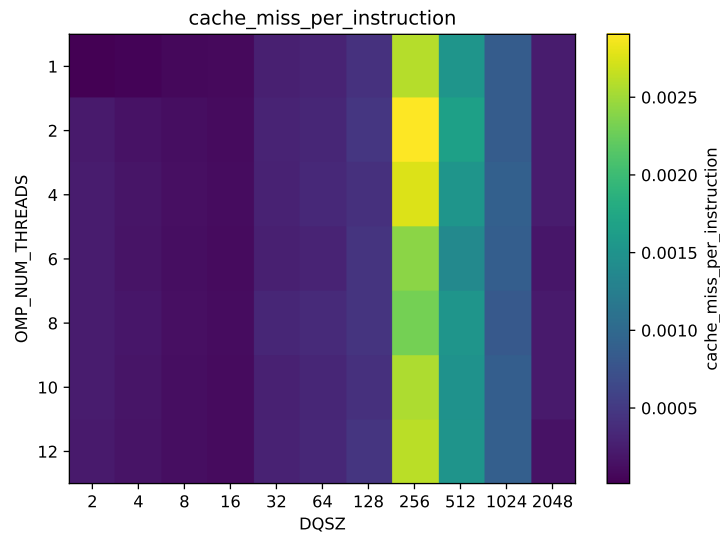## 3.2 Heatmap 2: Speedup (vs same-$DQSZ$, single-thread)

Interpretation:

- Speedup is highest for small base-case sizes (8 and 16).
- Large $DQSZ$ values generate too few tasks, limiting parallelism.
- Small $DQSZ$ values enable enough fine-grained tasks to keep all threads busy.

This graph explains that for some DQSZ we take much more advantage of parallelism, but doesn't prove that global execution time is better (because of memory bandwidth limitations).

3.3 Heatmap 3: Cache Misses per Instruction

Interpretation :

- Cache misses per instruction are minimal for $DQSZ = 8$ and $DQSZ = 16$.
- Their working sets fit inside caches (L1/L2 ?)

- Larger $DQSZ$ values overflow L1/L2 and drastically increase miss rate.

We find a sweet spot of DQSZ=16 for cache optimization, probably related to the size of the L1 cache compared to the DQSZ size in memory.

4   Conclusion